```
■ tiiuae/falcon-refinedweb
■ Viewer • Updated Jun 20, 2023 • ± 10.1k • ♡ 696

■ togethercomputer/RedPajama-Data-1T
■ Viewer • Updated Jun 30, 2023 • ± 3.91k • ♡ 976

■ bigcode/starcoderdata
Updated May 16, 2023 • ± 3.85k • ♡ 301

■ Spaces using openlm-research/open_llama_3b_v2 23

▼ HuggingFaceH4/open_llm_leaderboard
■ Sharathhebbar24/One-stop-for-Open-source-models
▼ gsalvinay/open_llm_leaderboard
▼ duagarcia/open_pt_llm_leaderboard
● GTBench/GTBench
+ 18 Spaces
```

Edit model card

OpenLLaMA: An Open Reproduction of LLaMA

TL;DR: we are releasing our public preview of OpenLLaMA, a permissively licensed open source reproduction of Meta Al's LLaMA. We are releasing a series of 3B, 7B and 13B models trained on different data mixtures. Our model weights can serve as the drop in replacement of LLaMA in existing implementations.

In this repo, we present a permissively licensed open source reproduction of Meta AI's <u>LLaMA</u> large language model. We are releasing a series of 3B, 7B and 13B models trained on 1T tokens. We provide PyTorch and JAX weights of pre-trained OpenLLaMA models, as well as evaluation results and comparison against the original LLaMA models. The v2 model is better than the old v1 model trained on a different data mixture. Please see the <u>project homepage of OpenLLaMA</u> for more details.

Weights Release, License and Usage

We release the weights in two formats: an EasyLM format to be use with our <u>EasyLM</u> framework, and a PyTorch format to be used with the <u>Hugging Face transformers</u> library. Both our training framework EasyLM and the checkpoint weights are licensed permissively under the Apache 2.0 license.

Loading the Weights with Hugging Face Transformers

Preview checkpoints can be directly loaded from Hugging Face Hub. Please note that it is advised to avoid using the Hugging Face fast tokenizer for now, as we've observed that the auto-converted fast tokenizer sometimes gives incorrect tokenizations. This can be achieved by directly using the LlamaTokenizer class, or passing in the use_fast=False option for the AutoTokenizer class. See the following example for usage.

```
import torch
from transformers import LlamaTokenizer, LlamaForCausalLM
## v2 models
model_path = 'openlm-research/open_llama_3b_v2'
# model_path = 'openlm-research/open_llama_7b_v2'
## v1 models
# model_path = 'openlm-research/open_llama_3b'
# model_path = 'openlm-research/open_llama_7b'
# model_path = 'openlm-research/open_llama_13b'
tokenizer = LlamaTokenizer.from_pretrained(model_path)
model = LlamaForCausalLM.from_pretrained(
    model_path, torch_dtype=torch.float16, device_map='auto',
)
prompt = 'Q: What is the largest animal?\nA:'
input_ids = tokenizer(prompt, return_tensors="pt").input_ids
generation_output = model.generate(
    input_ids=input_ids, max_new_tokens=32
```

```
print(tokenizer.decode(generation_output[0]))
```

For more advanced usage, please follow the <u>transformers LLaMA documentation</u>.

Evaluating with LM-Eval-Harness

The model can be evaluated with lm-eval-harness. However, due to the aforementioned tokenizer issue, we need to avoid using the fast tokenizer to obtain the correct results. This can be achieved by passing in use_fast=False to this can be achieved by passing in use_fast=False to this part of lm-eval-harness, as shown in the example below:

```
tokenizer = self.AUTO_TOKENIZER_CLASS.from_pretrained(
    pretrained if tokenizer is None else tokenizer,
    revision=revision + ("/" + subfolder if subfolder is not None else
    use_fast=False
)
```

Loading the Weights with EasyLM

For using the weights in our EasyLM framework, please refer to the <u>LLaMA</u> documentation of <u>EasyLM</u>. Note that unlike the original LLaMA model, our OpenLLaMA tokenizer and weights are trained completely from scratch so it is no longer needed to obtain the original LLaMA tokenizer and weights.

O Dataset and Training

The v1 models are trained on the <u>RedPajama dataset</u>. The v2 models are trained on a mixture of the <u>Falcon refined-web dataset</u>, the <u>StarCoder dataset</u> and the wikipedia, arxiv, book and stackexchange part of the <u>RedPajama dataset</u>. We follow the exactly same preprocessing steps and training hyperparameters as the original LLaMA paper, including model architecture, context length, training steps, learning rate schedule, and optimizer. The only difference between our setting and the original one is the

dataset used: OpenLLaMA employs open datasets rather than the one utilized by the original LLaMA.

We train the models on cloud TPU-v4s using <u>EasyLM</u>, a JAX based training pipeline we developed for training and fine-tuning large language models. We employ a combination of normal data parallelism and fully sharded data parallelism (also know as ZeRO stage 3) to balance the training throughput and memory usage. Overall we reach a throughput of over 2200 tokens / second / TPU-v4 chip for our 7B model.

⊘ Evaluation

We evaluated OpenLLaMA on a wide range of tasks using <u>lm-evaluation-harness</u>. The LLaMA results are generated by running the original LLaMA model on the same evaluation metrics. We note that our results for the LLaMA model differ slightly from the original LLaMA paper, which we believe is a result of different evaluation protocols. Similar differences have been reported in <u>this issue of lm-evaluation-harness</u>. Additionally, we present the results of GPT-J, a 6B parameter model trained on the <u>Pile</u> dataset by <u>EleutherAl</u>.

The original LLaMA model was trained for 1 trillion tokens and GPT-J was trained for 500 billion tokens. We present the results in the table below. OpenLLaMA exhibits comparable performance to the original LLaMA and GPT-J across a majority of tasks, and outperforms them in some tasks.

	GPT						
	-J	LLaMA	LLaMA	OpenLLaMA	OpenLLaMA	OpenLLaMA	(
Task/Metric	6B	7B	13B	3Bv2	7Bv2	3B	_
anli_r1/acc	0.32	0.35	0.35	0.33	0.34	0.33	(
anli_r2/acc	0.34	0.34	0.36	0.36	0.35	0.32	(
anli_r3/acc	0.35	0.37	0.39	0.38	0.39	0.35	(
arc_challenge/acc	0.34	0.39	0.44	0.34	0.39	0.34	(

Task/Metric	GPT -J 6B	LLaMA 7B	LLaMA 13B	OpenLLaMA 3Bv2	OpenLLaMA 7Bv2	OpenLLaMA 3B
arc_challenge/acc_norm	0.37	0.41	0.44	0.36	0.41	0.37
arc_easy/acc	0.67	0.68	0.75	0.68	0.73	0.69
arc_easy/acc_norm	0.62	0.52	0.59	0.63	0.70	0.65
boolq/acc	0.66	0.75	0.71	0.66	0.72	0.68
hellaswag/acc	0.50	0.56	0.59	0.52	0.56	0.49
hellaswag/acc_norm	0.66	0.73	0.76	0.70	0.75	0.67
openbookqa/acc	0.29	0.29	0.31	0.26	0.30	0.27
openbookqa/acc_norm	0.38	0.41	0.42	0.38	0.41	0.40
piqa/acc	0.75	0.78	0.79	0.77	0.79	0.75
piqa/acc_norm	0.76	0.78	0.79	0.78	0.80	0.76
record/em	0.88	0.91	0.92	0.87	0.89	0.88
record/f1	0.89	0.91	0.92	0.88	0.89	0.89
rte/acc	0.54	0.56	0.69	0.55	0.57	0.58
truthfulqa_mc/mc1	0.20	0.21	0.25	0.22	0.23	0.22
truthfulqa_mc/mc2	0.36	0.34	0.40	0.35	0.35	0.35
wic/acc	0.50	0.50	0.50	0.50	0.50	0.48
winogrande/acc	0.64	0.68	0.70	0.63	0.66	0.62
Average	0.52	0.55	0.57	0.53	0.56	0.53

We removed the task CB and WSC from our benchmark, as our model performs suspiciously high on these two tasks. We hypothesize that there could be a benchmark

data contamination in the training set.

⊘ Contact

We would love to get feedback from the community. If you have any questions, please open an issue or contact us.

OpenLLaMA is developed by: <u>Xinyang Geng*</u> and <u>Hao Liu*</u> from Berkeley AI Research. *Equal Contribution

Acknowledgment

We thank the <u>Google TPU Research Cloud</u> program for providing part of the computation resources. We'd like to specially thank Jonathan Caton from TPU Research Cloud for helping us organizing compute resources, Rafi Witten from the Google Cloud team and James Bradbury from the Google JAX team for helping us optimizing our training throughput. We'd also want to thank Charlie Snell, Gautier Izacard, Eric Wallace, Lianmin Zheng and our user community for the discussions and feedback.

The OpenLLaMA 13B v1 model is trained in collaboration with <u>Stability AI</u>, and we thank Stability AI for providing the computation resources. We'd like to especially thank David Ha and Shivanshu Purohit for the coordinating the logistics and providing engineering support.

Reference

If you found OpenLLaMA useful in your research or applications, please cite using the following BibTeX:

```
@software{openlm2023openllama,
  author = {Geng, Xinyang and Liu, Hao},
  title = {OpenLLaMA: An Open Reproduction of LLaMA},
  month = May,
```

```
year = 2023,
url = {https://github.com/openlm-research/open_llama}
}

@software{together2023redpajama,
author = {Together Computer},
title = {RedPajama-Data: An Open Source Recipe to Reproduce LLaMA tramonth = April,
year = 2023,
url = {https://github.com/togethercomputer/RedPajama-Data}
}
```

```
@article{touvron2023llama,
  title={Llama: Open and efficient foundation language models},
  author={Touvron, Hugo and Lavril, Thibaut and Izacard, Gautier and Majournal={arXiv preprint arXiv:2302.13971},
  year={2023}
}
```



Company

TOS

Privacy

About

Jobs

Website

Models
Datasets
Spaces

Pricing

Docs

© Hugging Face